



Qwen3-8B accuracy across decoding and training seeds

DeepMath-64 training → AMC23/AIME24/AIME25 (143 problems) · thinking mode · no LLM judge


a Decoding-seed sensitivity
fixed seed-0 training run · 3 decoding seeds



b Training-seed reproducibility
fixed decoding seed · reused seed-0 + 2 new H100 seeds



c Training time + checkpoint space
Qwen3-8B · two new H100 training seeds



Yellow = LGSD · Blue = OPSD · point estimates; no error bars