


Does an answer-free prompt paraphrase change the distillation target?

Same student prefix and token; only the neutral / terse / verbose privileged prompt changes.

a

How one point is built: style token “Therefore”



Descriptive same-prefix case study: 3 DeepMath queries $\times$ 3 answer-free prompt wrappers; token points are paired observations, not independent query replicates. Exact zeros are displayed at a visual floor.

b

Every token is paired before vs. after projection (24,140 tokens)

