



Qwen3-8B training process · 64 trajectories, thousands of chunk-level steps

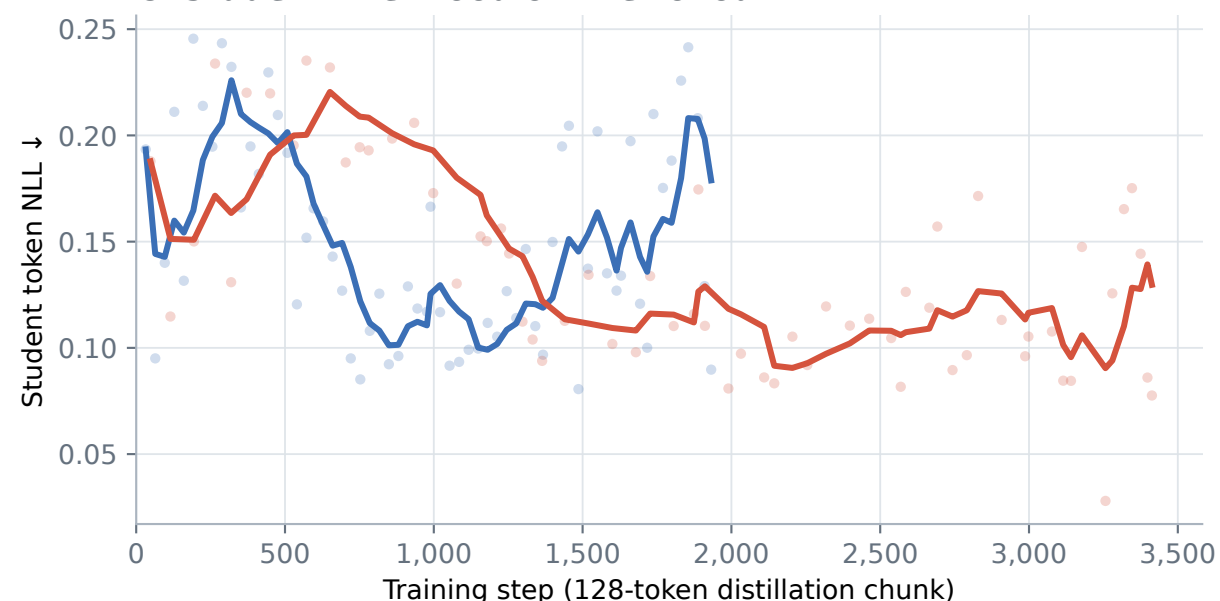
A Distillation loss



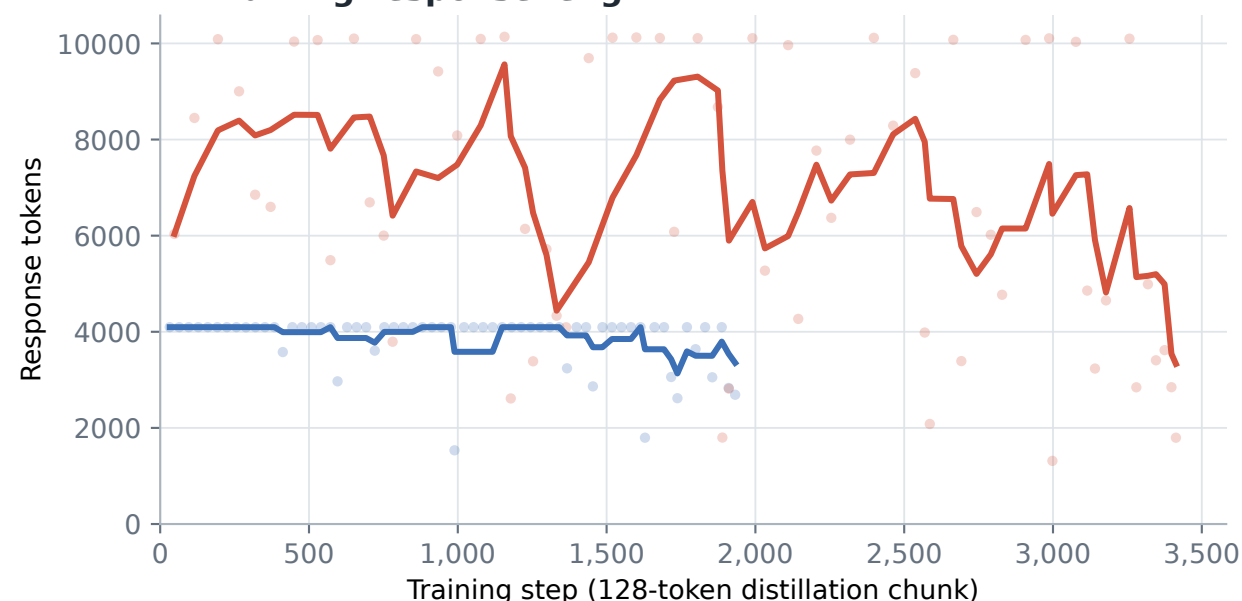
B Teacher-student divergence



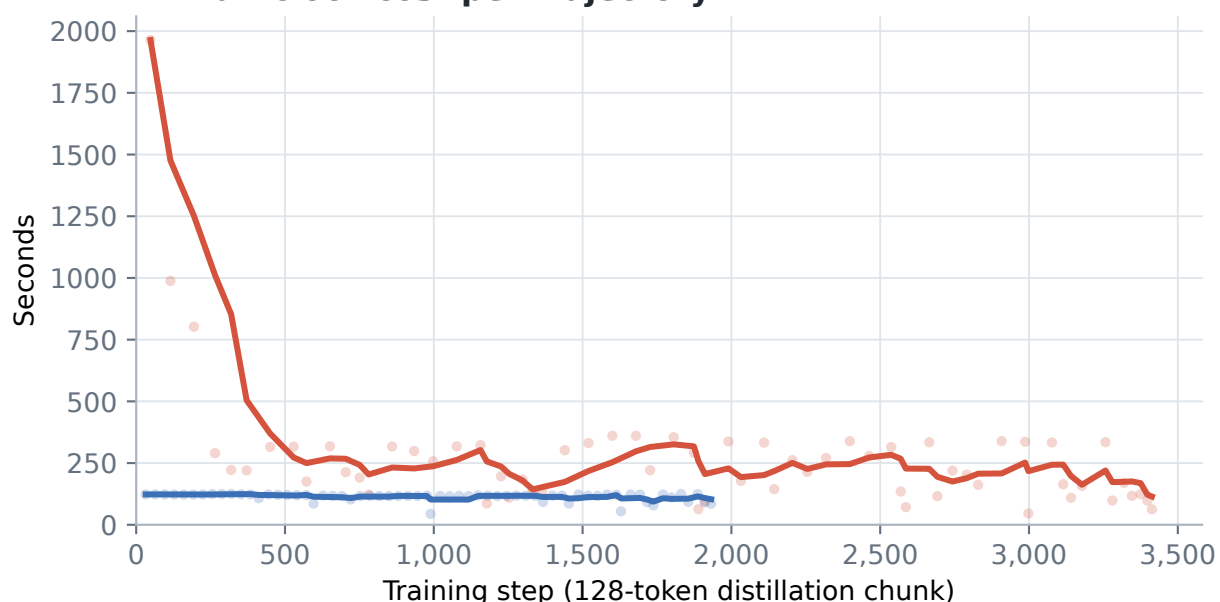
C Student likelihood on the rollout



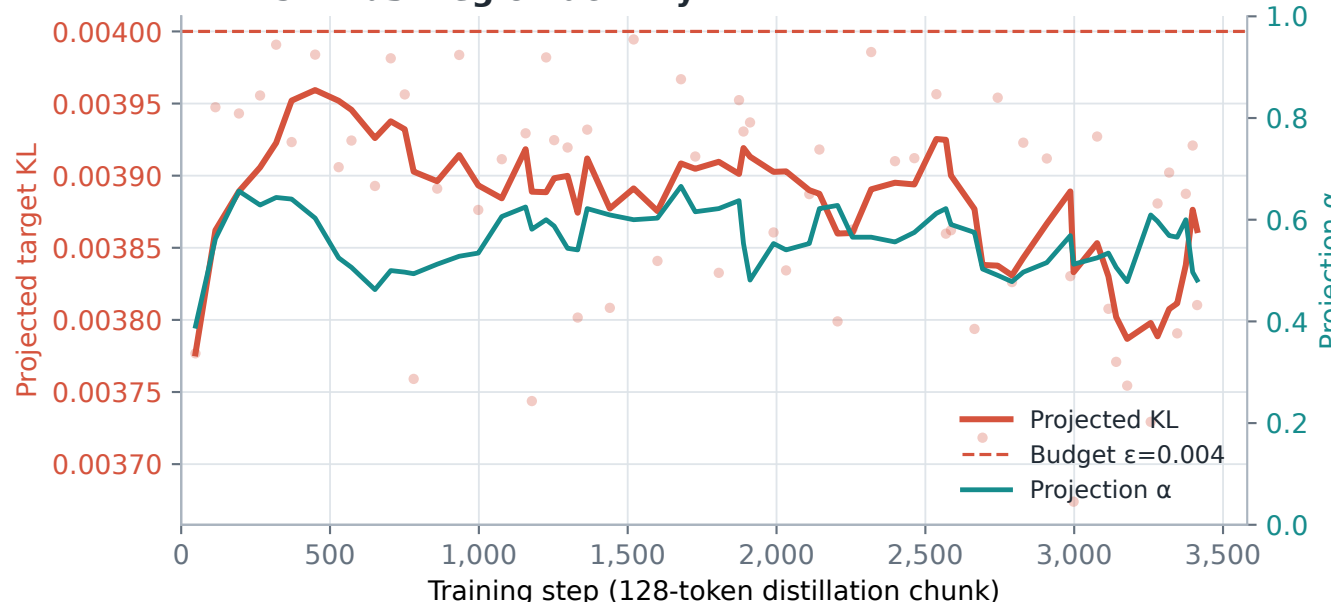
D Training-response length



E Wall-clock cost per trajectory



F TRSD trust-region activity



Privilege-SD 64: 1,932 steps, TRSD 64: 3,413 steps. Each point is a trajectory aggregate at its exact cumulative chunk boundary.