



A Context preferences become training targets



Arrows show implemented dependencies; harmful feedback is a hypothesis about this loop.

B Shrinking a harmful preference

Synthetic example · not GPT-OSS measurements



C Local control preserves gradient direction

$$\nabla L_{\text{LGSD}}(\theta_k) = \alpha_k \nabla L_{\text{raw}}(\theta_k)$$

At an identical state: detached exponential target, unclipped reverse KL, one shared α .

A positive α preserves the local gradient direction. It supplies no test of whether that direction improves task correctness.

Established: objective and parameter-sharing loop.
Unresolved: the harmful GPT-OSS token preference.